

# Single Core Design Space Exploration

ECE 568: Advanced Microprocessor Architecture, Spring 2016

Anthony Gallotta  
agallo4@uic.edu

## ABSTRACT

This paper seeks to find the best single core superscalar architecture for a given benchmark application by experimenting with different architectures for branch prediction, memory systems, functional units, data path, and other areas. The metrics used to define the "best" architecture are *performance* as instructions per cycle (IPC) and *performance-energy* as the energy delay product (EDP).

## 1. INTRODUCTION

This study evaluates a single core superscalar architecture using the SimpleScalar simulator suite, seeking the best configuration for the **eeg** benchmark application. A base architecture was given, and design modifications were then made in four groups, branch prediction, memory system, functional units, and data path. In the first phase of design space exploration, tuning one parameter at a time. In the second stage, closely related parameters within the same groups are tuned together. Finally, the best configurations across groups are combined to formulate the best found architecture. For each configuration we observe performance (IPC) and performance-energy (EDP), calculated as  $(CPI * energy/cycle) * CPI$ . The "best" configuration seeks to maximize performance, and minimize performance-energy.

## 2. SIMULATION

Simulations of the **eeg** application benchmark were performed using **sim-outorder**. The base configuration had performance of 1.5226 IPC, and a performance-energy product of 219.68. While many parameter changes were tested, unless otherwise noted, all graphs that follow show only configurations that outperformed the base configuration in at least one of these metrics. The most desirable configurations will lie in the lower right side of the graphs that follow.

To gain some direction in performance tuning, the instruction profile for the benchmark was first evaluated using **sim-profile**. As shown in table 1, the benchmark is primarily

Table 1: Instruction Profile for eeg

Type	Count	Percent
load	171,182,879	20.30
store	84,069,724	9.97
uncond branch	60,740,822	7.20
cond branch	59,922,458	7.11
int computation	335,861,146	39.84
fp computation	131,306,086	15.57
trap	1,035	0.00

an integer program. Given this instruction breakdown, most of the performance tuning was focused on branch prediction and the memory system.

### 2.1 Branch Prediction

The first branch prediction change attempted was to use static prediction. This strategy resulted in much worse performance by both metrics and was quickly abandoned. Next, some tuning was made with the default bimodal branch predictor, which uses a simple strategy of picking the most common direction. Sticking with this scheme, the branch target buffer (BTB) size and associativity was adjusted. In figure 2.1, the points labeled as **bp-btb-<sets>-<associativity>** indicate the results of these adjustments. The BTB configuration of 128 sets with 2 way associativity is able to achieve the same performance as the base configuration, with a much lower EDP.

Two level branch predictors use a combination of the last  $k$  historical branch outcomes, as well as the behavior for a specific pattern of previous branches [1]. While many of the 2 level configurations did not perform the base configuration, the 2 level Global Address (GA) branch predictor with a 4 wide history (bp-2level-hist-4) did achieve a better IPC. Combining these configuration changes (bp-comb-128-2) by keeping the best found BTB size and using a combined bimodal and 2-level branch predictor gives about a 5% performance-energy improvement over the base configuration.

A perfect branch prediction scheme is plotted for comparison. Surprisingly, the perfect branch predictor achieves a lower IPC than many of the experimental configurations. The simulator documentation does not describe in detail how the perfect branch predictor operates, so it may be that they are modeling a higher latency for the perfect predictor.

### 2.2 Memory System

The simulator has a 2 level cache hierarchy, and we ex-

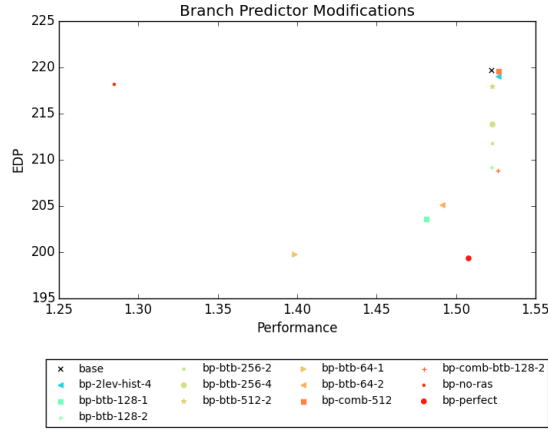


Figure 1: Branch Predictor Variation results




Figure 3: Modifications for L1 instruction cache




Figure 2: Modifications for L1 data cache

periment with different sizes, associativity, block sizes, and replacement policies at each level, and for both instruction and data caches.

### 2.2.1 L1 Cache

As figure 2 shows, increasing the size of the L1 data cache (sets) did not result in significant IPC improvements relative to their increased power consumption. Increasing the block size to 64 while keeping the total L1 size the same (bsize-64) yields the best result in terms of both of our architecture goals.

Figure 3 shows similar results for the instruction cache - a block size of 64 performs best.

### 2.2.2 L2 Cache

At the L2 cache, we see once again from figures 4, 5, and 6 that a unified cache with a decreased block size performs best.

## 3. CONCLUSIONS

This paragraph will end the body of this sample document. Remember that you might still have Acknowledgments or Appendices; brief samples of these follow. There is

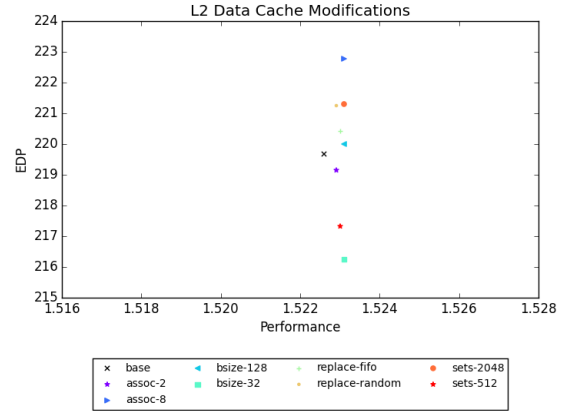


Figure 4: Modifications for L2 data cache





Figure 5: Modifications for L2 instruction cache



**Figure 6: Unified L2 cache**

still the Bibliography to deal with; and we will make a disclaimer about that here: with the exception of the reference to the L<sup>A</sup>T<sub>E</sub>X book, the citations in this paper are to articles which have nothing to do with the present subject and are used as examples only.

## 4. REFERENCES

- [1] T.-Y. Yeh and Y. N. Patt. A comparison of dynamic branch predictors that use two levels of branch history. In *ACM SIGARCH Computer Architecture News*, volume 21, pages 257–266. ACM, 1993.